

## Transformer-Based SLAM Learning Checklist (PyTorch + CUDA)

---

### Phase 1 – PyTorch Foundations (Weeks 1–3)

**Goal:** Master PyTorch, data loading, and attention basics.

- Complete [PyTorch Tutorials](#).
- Implement Linear Regression & CNN manually.
- Inspect graphs with `torchsummary` or `torchviz`.
- Write custom `Dataset` & `DataLoader` using [Data Loading Tutorial](#).
- Implement toy Vision Transformer (ViT) using [ViT PyTorch \(lucidrains\)](#).
- Visualize attention maps.

 Deliverables: - `train_vit_toy.py` - Attention visualization notebook

---

### Phase 2 – Transformer Feature Matching (Weeks 4–6)

**Goal:** Implement LoFTR-style matching on paired images.

- Read [LoFTR Paper](#).
- Clone & study [Official LoFTR Repo](#).
- Implement coarse feature matching (shared CNN + cross-attention).
- Train on KITTI or TUM-RGBD pairs.
- Use [Kornia Matching Tutorial](#).
- Train using mixed precision ([AMP Recipe](#)).

 Deliverables: - `loftr_simplified.py` - Heatmaps + validation loss plots

---

### Phase 3 – Pose Estimation & CUDA (Weeks 7–9)

**Goal:** Estimate pose and optimize model using CUDA.

- Extract top-N matches from Transformer output.
- Estimate pose using `cv2.findEssentialMat` + `cv2.recoverPose` ([OpenCV Docs](#)).
- Evaluate translation/rotation errors.
- Profile using `torch.profiler` & Nsight Systems.
- Implement one custom CUDA kernel ([Tutorial](#)).
- Run ablations (embedding dim, attention heads, etc.).

 Deliverables: - `pose_from_correspondences.py` - Pose trajectory + GPU profile logs

---

## Phase 4 – Planning Extension (Weeks 10–12)

**Goal:** Extend Transformer pipeline to planning or trajectory prediction.

- Replace image inputs with state sequences `[x, y, θ, v, ω]`.
- Train Decision Transformer or Diffusion Planner variant.
- Study [Decision Transformer Paper](#).
- Use [Decision Transformer Repo](#).
- Try [Diffusion Policy Repo](#).
- Compare perception vs. planning data format in a notebook.

 Deliverables: - `transformer_planner.py` - Comparative visualization notebook

---

## Final Outputs

- Fully working Transformer SLAM pipeline
  - CUDA kernel + profiling results
  - Extended planning pipeline
  - Portfolio-ready documentation
- 

## Core Tools Summary

Category	Tools
Deep Learning	PyTorch, torchvision, Kornia
CUDA & Profiling	torch.profiler, Nsight Systems
Vision & Geometry	OpenCV, numpy
Datasets	KITTI, TUM-RGBD
Visualization	matplotlib, evo
Environment	Ubuntu + Conda (GPU RTX 3060+ recommended)

---

This checklist ensures week-by-week tangible progress from PyTorch mastery → Transformer SLAM → CUDA optimization → Transformer-based planning.